一 基于信任 Gateway Outbound 的流量模型

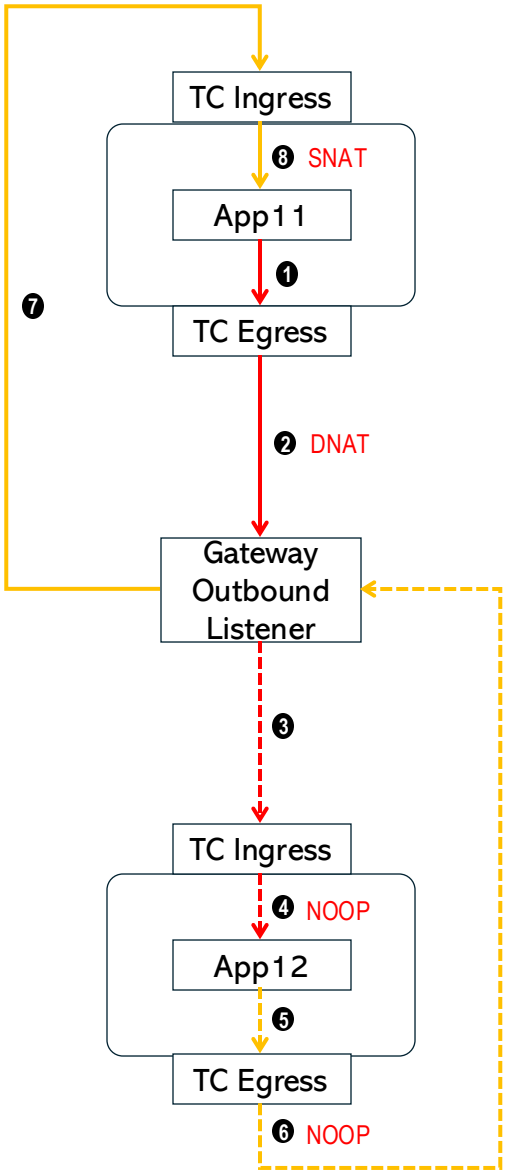

Mesh内互访

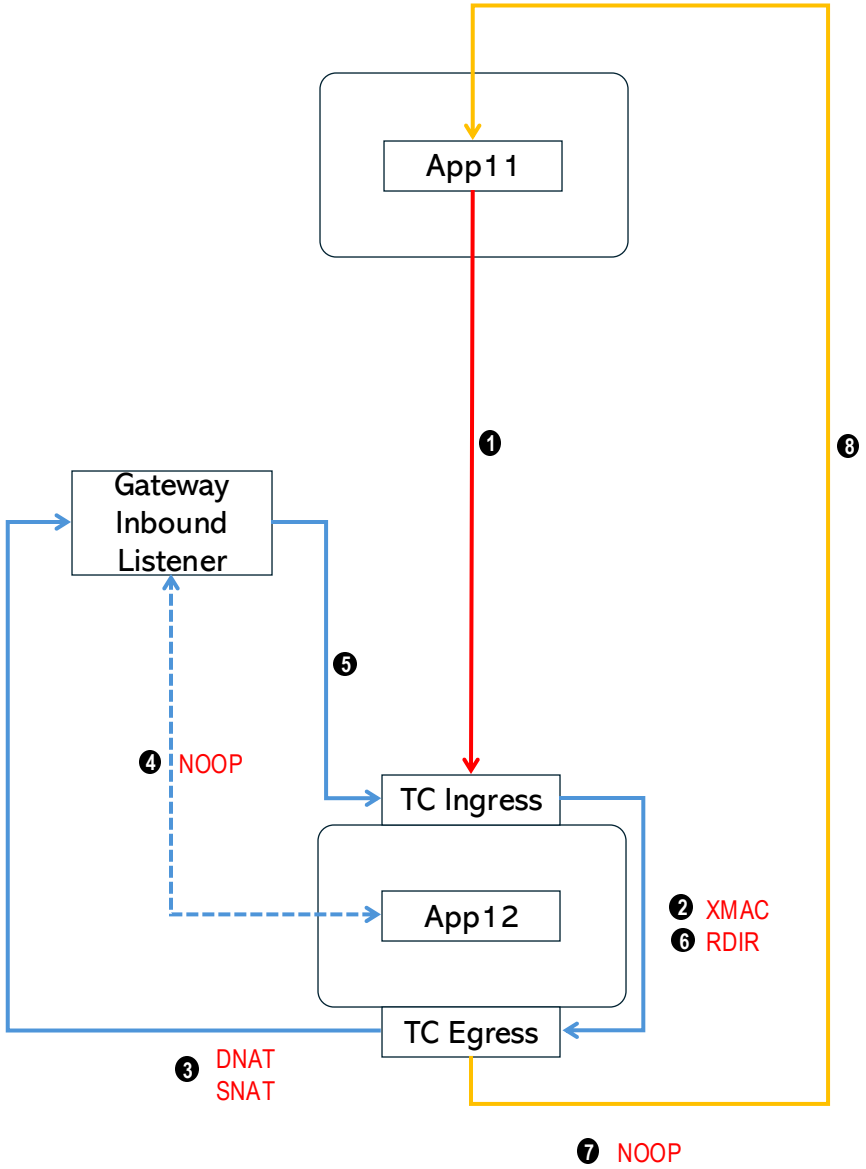

Mesh外访问Mesh内

App进/出的流量只会被拦截给同 Node 上的 FGW Listeners

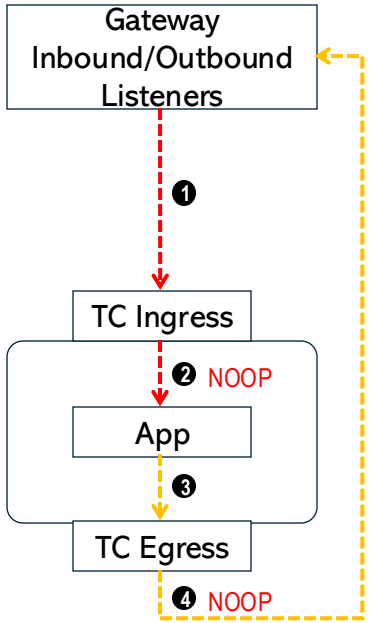

Gateway访问Mesh内

二 Gateway Outbound 信任与否的流量模型比较

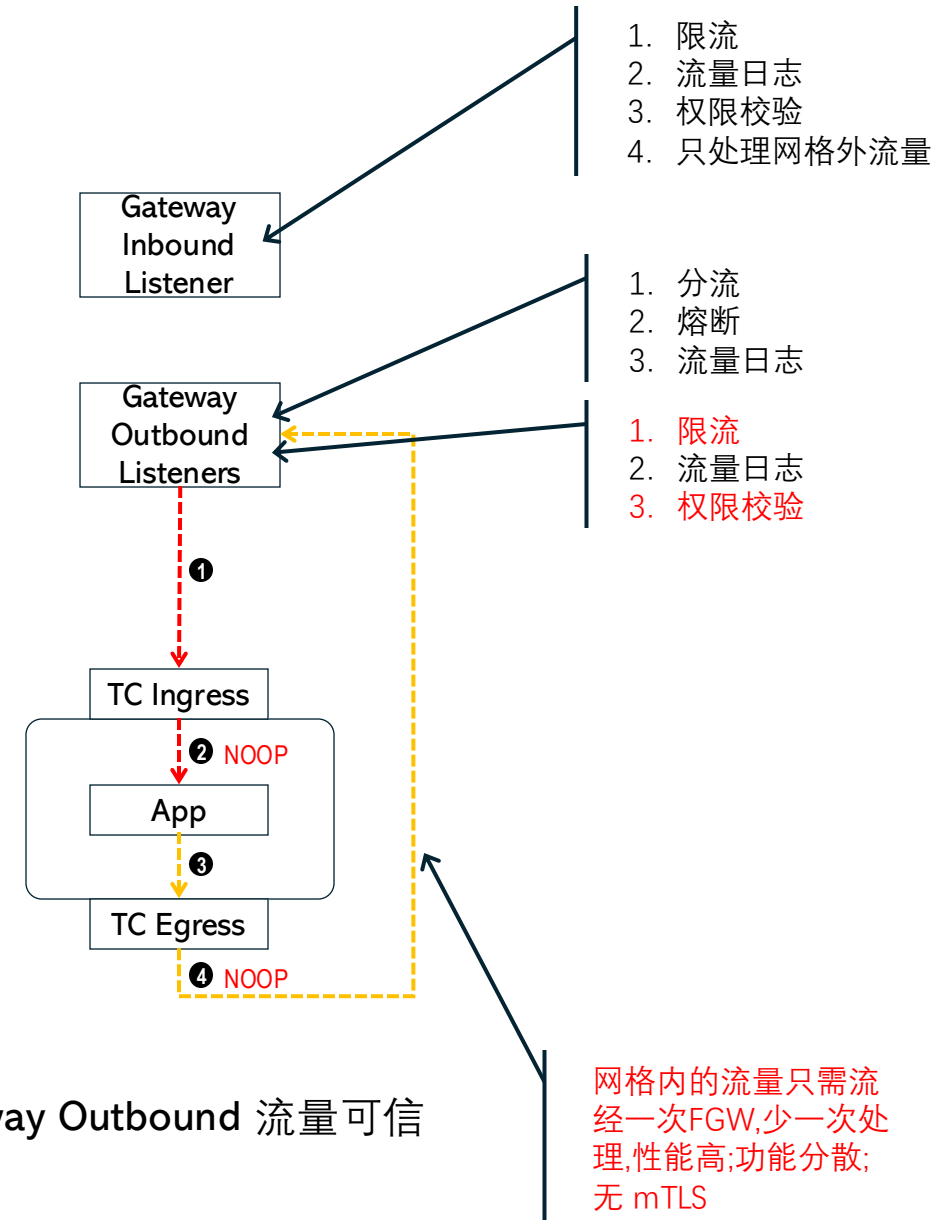

Gateway Outbound 流量可信

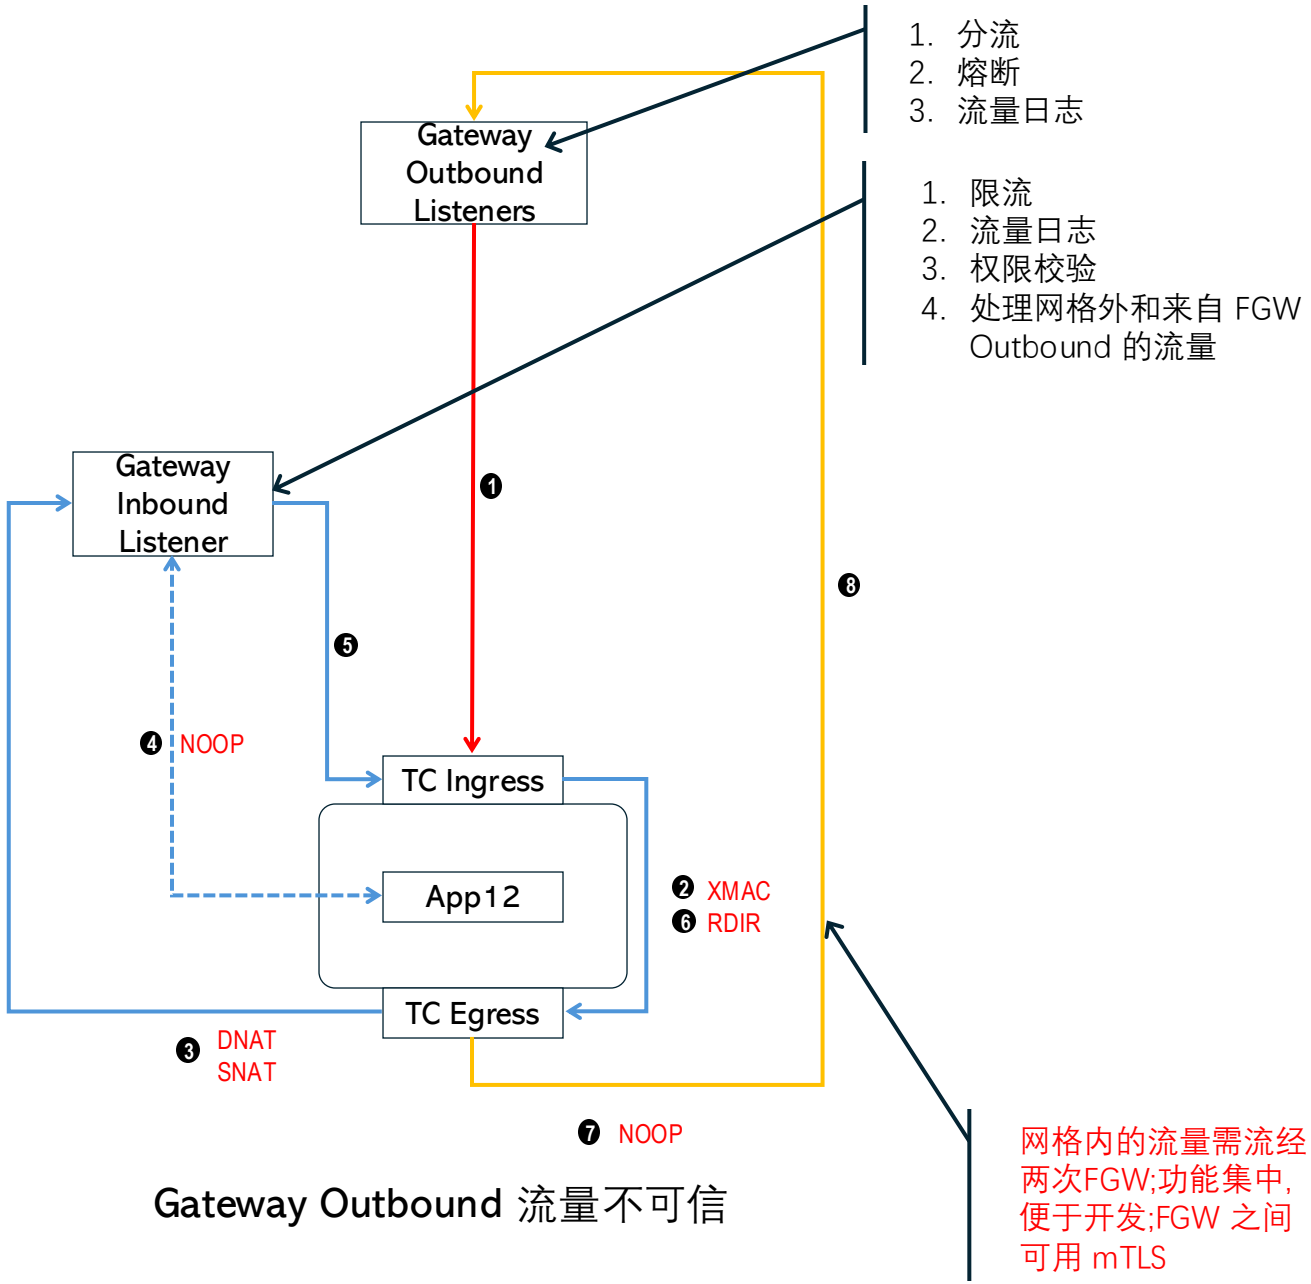

Gateway Outbound 流量不可信

三 Mesh 各控制层服务部署示意图

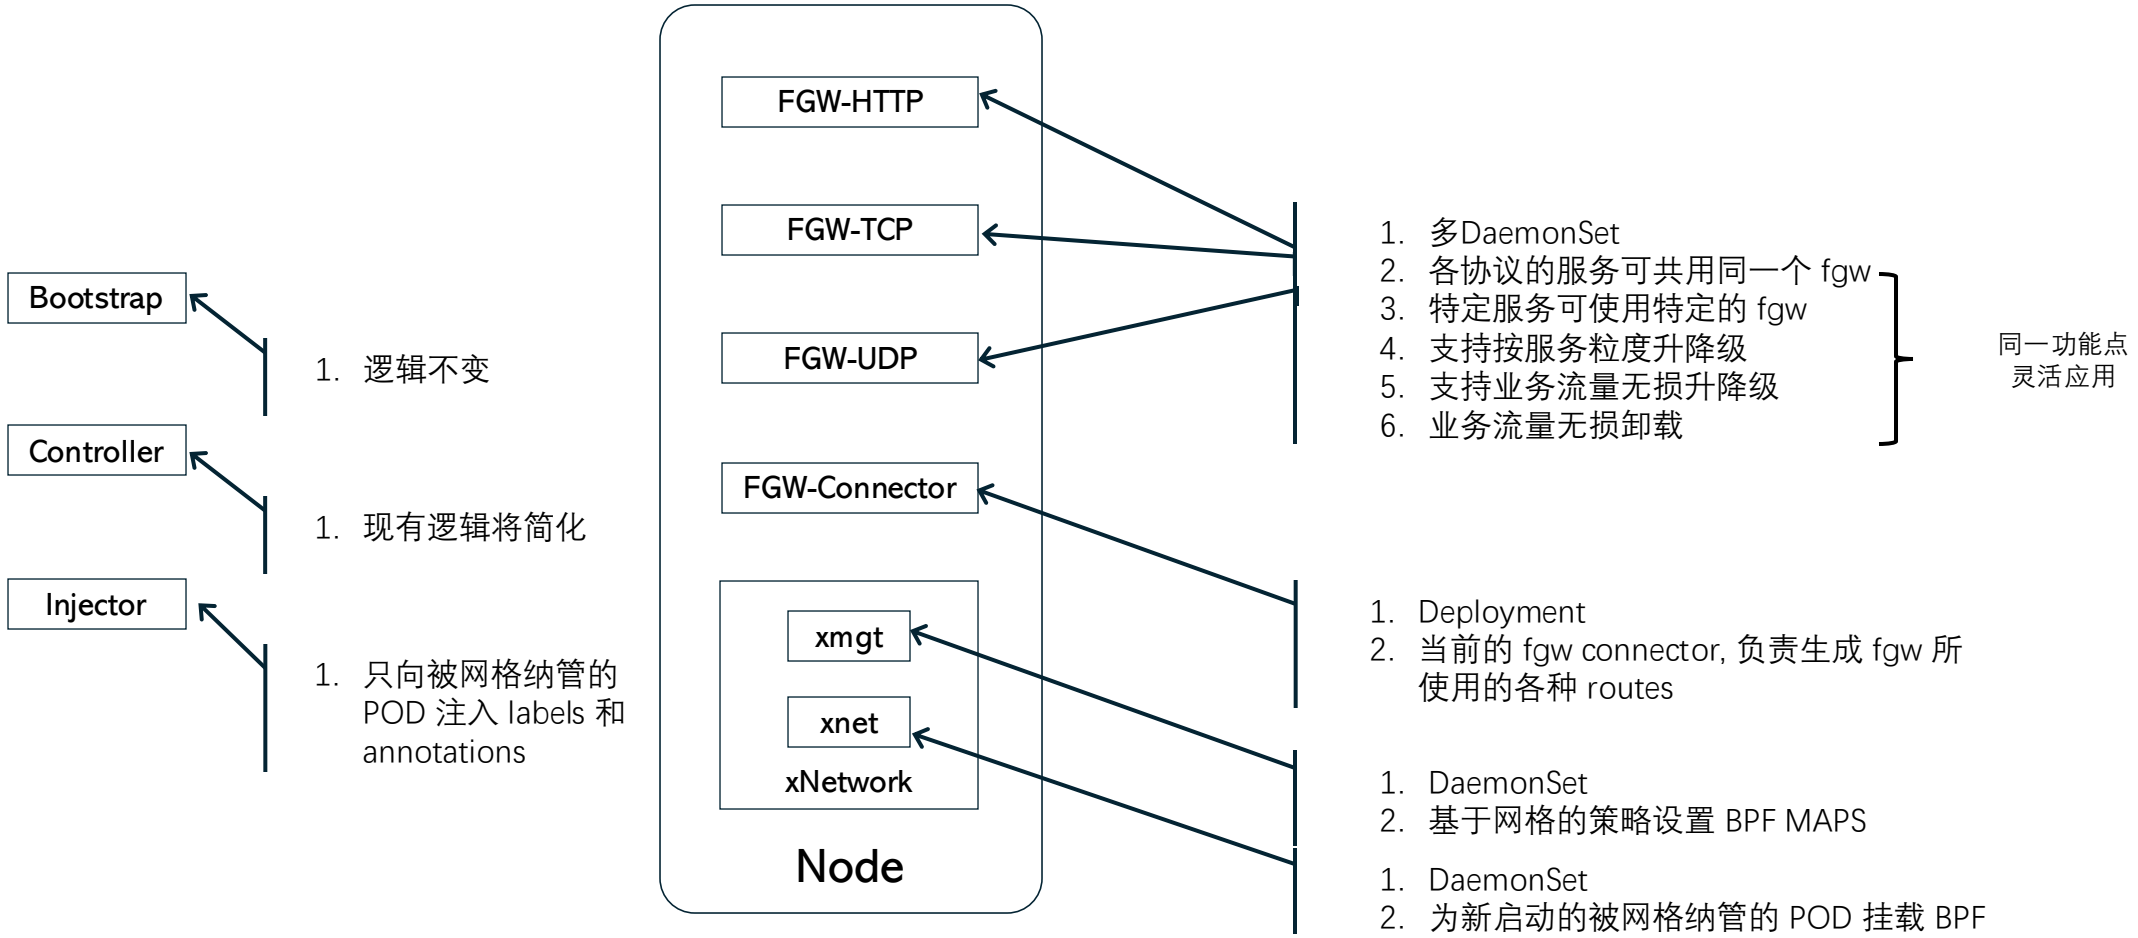

四 eBPF流量处理逻辑流程

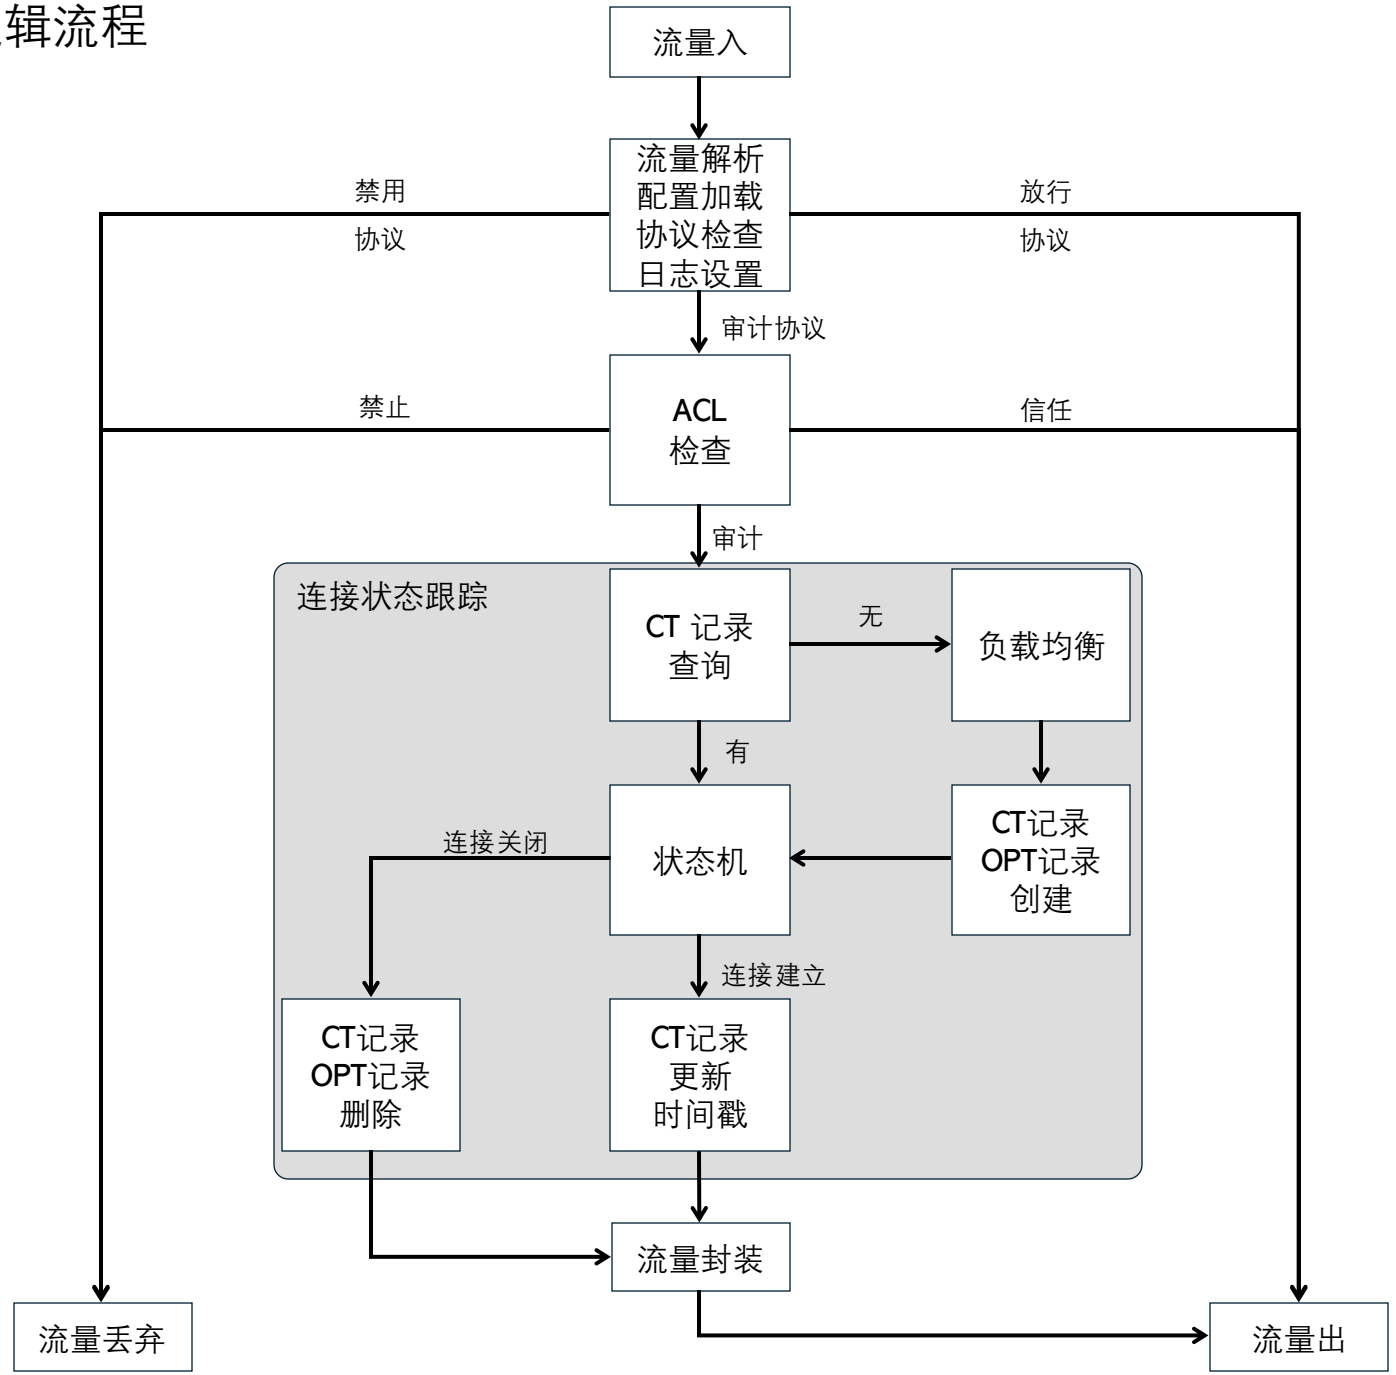

TC 的 Ingress 和 Egress 相同的处理流程, 处理的方向相反

五 主要的 eBPF maps


```
[
  {
    "key": {
      "daddr": "0.0.0.0",
      "dport": 0,
      "proto": "IPPROTO_TCP",
      "v6": false,
      "tc_dir": "TC_DIR_IGR"
    },
    "value": {
      "ep_sel": 0,
      "ep_cnt": 1,
      "eps": [
        {
          "raddr": "192.168.226.22",
          "rport": 15003,
          "inactive": false
        }
      ]
    }
  },
  {
    "key": {
      "daddr": "0.0.0.0",
      "dport": 0,
      "proto": "IPPROTO_TCP",
      "v6": false,
      "tc_dir": "TC_DIR_EGR"
    },
    "value": {
      "ep_sel": 0,
      "ep_cnt": 1,
      "eps": [
        {
          "raddr": "192.168.226.22",
          "rport": 15001,
          "inactive": false
        }
      ]
    }
  }
]
```


fsm_xnat


```
[
  {
    "key": {
      "daddr": "10.0.0.2",
      "saddr": "192.168.226.22",
      "dport": 43550,
      "sport": 15001,
      "proto": "IPPROTO_TCP",
      "v6": false
    },
    "value": {
      "flow_dir": "FLOW_DIR_S2C",
      "to": 0,
      "ts": 8818218019373,
      "nfs": {
        "TC_DIR_IGR": "NF_ALLOW NF_XNAT",
        "TC_DIR_EGR": "NF_DENY"
      },
      "do_trans": false,
      "xnat": {
        "xaddr": "0.0.2.0",
        "raddr": "0.0.2.0",
        "xport": 37034,
        "rport": 7681,
        "fin": false
      },
      "trans": {
        "tcp": {
          "state": "TCP_STATE_CLOSED",
          "fin_dir": "",
          "conns": {
            "FLOW_DIR_C2S": {
              "seq": 2678153596,
              "prev_ack_seq": 346257579,
              "prev_seq": 0,
              "init_acks": 0
            },
            "FLOW_DIR_S2C": {
              "seq": 0,
              "prev_ack_seq": 0,
              "prev_seq": 0,
              "init_acks": 0
            }
          }
        },
        "udp": {
          "_state": 0,
          "_pkts_seen": 0,
          "_rpkts_seen": 0,
          "_fin_dir": "FLOW_DIR_C2S"
        }
      }
    }
  },
```



```
  {
    "key": {
      "daddr": "20.0.0.2",
      "saddr": "10.0.0.2",
      "dport": 8080,
      "sport": 43550,
      "proto": "IPPROTO_TCP",
      "v6": false
    },
    "value": {
      "flow_dir": "FLOW_DIR_C2S",
      "to": 0,
      "ts": 8818218019373,
      "nfs": {
        "TC_DIR_IGR": "NF_DENY",
        "TC_DIR_EGR": "NF_ALLOW NF_XNAT"
      },
      "do_trans": false,
      "xnat": {
        "xaddr": "0.0.2.0",
        "raddr": "168.226.22.0",
        "xport": 7738,
        "rport": 39169,
        "fin": false
      },
      "trans": {
        "tcp": {
          "state": "TCP_STATE_CLOSED",
          "fin_dir": "",
          "conns": {
            "FLOW_DIR_C2S": {
              "seq": 363034795,
              "prev_ack_seq": 2678153596,
              "prev_seq": 0,
              "init_acks": 41748
            },
            "FLOW_DIR_S2C": {
              "seq": 0,
              "prev_ack_seq": 0,
              "prev_seq": 0,
              "init_acks": 0
            }
          }
        },
        "udp": {
          "_state": 0,
          "_pkts_seen": 0,
          "_rpkts_seen": 0,
          "_fin_dir": "FLOW_DIR_C2S"
        }
      }
    }
  }
]
```


fsm_xflow


```
[
  {
    "key": {
      "saddr": "10.0.0.2",
      "sport": 43550,
      "proto": "IPPROTO_TCP"
    },
    "value": {
      "daddr": "20.0.0.2",
      "saddr": "10.0.0.2",
      "sport": 43550,
      "dport": 8080,
      "proto": "IPPROTO_TCP",
      "v6": false
    }
  }
]
```


fsm_xopt


```
[
  {
    "key": {
      "addr": "10.0.0.1",
      "port": 0,
      "proto": "IPPROTO_TCP"
    },
    "value": {
      "acl": "ACL_TRUSTED"
    }
  },
  {
    "key": {
      "addr": "20.0.0.2",
      "port": 8080,
      "proto": "IPPROTO_TCP"
    },
    "value": {
      "acl": "ACL_TRUSTED"
    }
  }
]
```


fsm_xacl

五 主要的 eBPF maps

```
83: percpu_array  name fsm_cxpkt  flags 0x0
        key 4B  value 117B  max_entries 1  memlock 16384B
        btf_id 263
84: prog_array  name fsm_prog  flags 0x0
        key 4B  value 4B  max_entries 2  memlock 4096B
        owner_prog_type sched_cls  owner jited
        btf_id 263
85: hash  name fsm_xacl  flags 0x0
        key 19B  value 1B  max_entries 4096  memlock 98304B
        btf_id 263
86: hash  name fsm_xnat  flags 0x0
        key 21B  value 28B  max_entries 64  memlock 4096B
        btf_id 263
87: percpu_array  name fsm_cflop  flags 0x0
        key 4B  value 104B  max_entries 2  memlock 28672B
88: hash  name fsm_xflow  flags 0x0
        key 38B  value 104B  max_entries 1048576  memlock 150994944B
        btf_id 263
89: hash  name fsm_xopt  flags 0x1
        key 19B  value 38B  max_entries 1048576  memlock 67108864B
        btf_id 263
90: array  name fsm_xcfg  flags 0x0
        key 4B  value 8B  max_entries 1  memlock 4096B
        btf_id 263
91: hash  name fsm_trip  flags 0x1
        key 16B  value 2B  max_entries 16  memlock 4096B
        btf_id 263
92: hash  name fsm_trpt  flags 0x1
        key 2B  value 2B  max_entries 16  memlock 4096B
        btf_id 263
```


```
[
  {
    "flags": {
      "mask": 1111110000000000000,
      "ipv6_proto_deny_all": false,
      "ipv4_tcp_proto_deny_all": false,
      "ipv4_tcp_proto_allow_all": false,
      "ipv4_udp_proto_deny_all": false,
      "ipv4_udp_proto_allow_all": false,
      "ipv4_oth_proto_deny_all": false,
      "ipv4_tcp_nat_by_ip_port_on": false,
      "ipv4_tcp_nat_by_ip_on": false,
      "ipv4_tcp_nat_all_off": false,
      "ipv4_udp_nat_by_ip_port_on": false,
      "ipv4_udp_nat_by_ip_on": false,
      "ipv4_udp_nat_all_off": false,
      "ipv4_nat_orig_opt_on": false,
      "ipv4_acl_check_on": true,
      "ipv4_trace_hdr_on": true,
      "ipv4_trace_nat_on": true,
      "ipv4_trace_opt_on": true,
      "ipv4_trace_acl_on": true,
      "ipv4_trace_flow_on": true,
      "ipv4_trace_by_ip_on": false,
      "ipv4_trace_by_port_on": false
    }
  }
]
```


fsm_xcfg


```
[
  {
    "key": {
      "addr": "10.0.0.2"
    },
    "value": {
      "trace_tc_ingress_on": "true",
      "trace_tc_egress_on": "true"
    }
  }
]
```


fsm_trip


```
[
  {
    "key": {
      "port": 8080
    },
    "value": {
      "trace_tc_ingress_on": "true",
      "trace_tc_egress_on": "true"
    }
  }
]
```


fsm_trpt

六 访问控制策略

1. 优先按 IP+PORT查询, 其次是按 IP 查询
2. FGW 的 inbound 和 outbound 端口也要设置 ACL,策略为 AUDIT, 以便能做反向的 DNAT/SNAT

七 后续改进点

1. TC Ingress 前增加 XDP, 处理 ACL, 放行的流量不进入内核网络层, 直接转发出去
2. App 到 FGW的 Inbound 和 Outbound Listeners 的流量都是同节点流量, 适合加速处理; 需要较高的内核版本, 不适用一些国产 OS